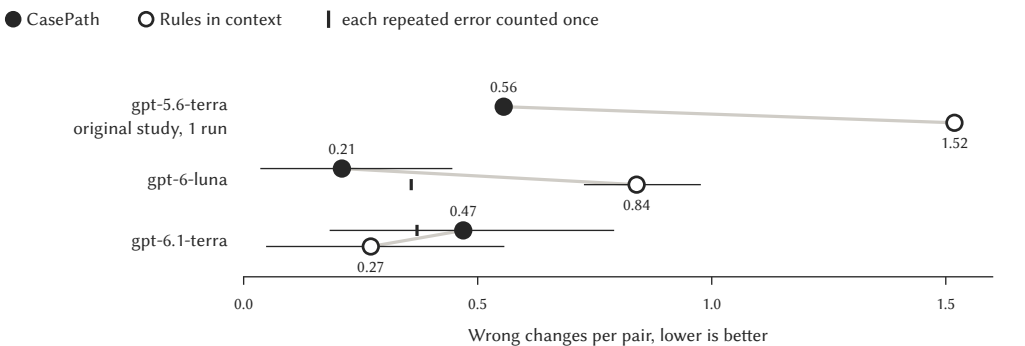
CasePath
Rules in context
each repeated error counted once
gpt-5.6-terra
original study, 1 run
0.56
1.52
gpt-6-luna
0.21
0.84
gpt-6.1-terra
0.47
0.27
0.0
0.5
1.0
1.5
Wrong changes per pair, lower is better