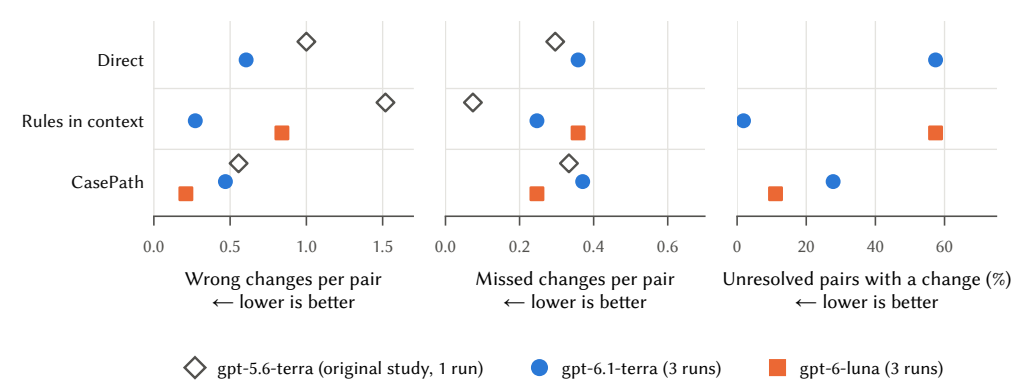
Direct
Rules in context
CasePath
0.0
0.5
1.0
1.5
Wrong changes per pair
← lower is better
0.0
0.2
0.4
0.6
Missed changes per pair
← lower is better
0
20
40
60
Unresolved pairs with a change (%)
← lower is better
gpt-5.6-terra (original study, 1 run)
gpt-6.1-terra (3 runs)
gpt-6-luna (3 runs)